

EVENT-DRIVEN VISUAL NAVIGATION FOR PRECISION PLANETARY LANDING IN EXTREME ENVIRONMENTS

A Neuromorphic Vision Approach Using
Spiking Neural Networks

A THESIS

Submitted in Partial Fulfillment of the Requirements
for the Degree of

DOCTOR OF PHILOSOPHY

in

Aerospace Engineering

Department of Aerospace Engineering
School of Engineering and Applied Sciences

February 2026

© 2026

All Rights Reserved

ABSTRACT

Autonomous precision landing on planetary bodies represents one of the most challenging problems in spacecraft navigation, requiring robust perception systems capable of operating in extreme environmental conditions while meeting stringent real-time computational constraints. Traditional frame-based visual sensors suffer from fundamental limitations including motion blur, limited dynamic range, and high bandwidth requirements that compromise their effectiveness during the critical powered descent phase of planetary landing missions.

This dissertation presents a comprehensive investigation into event-driven visual navigation using neuromorphic sensors and bio-inspired processing for spacecraft precision landing applications. The research addresses the fundamental question: *Can neuromorphic vision systems provide robust, low-latency visual navigation for autonomous planetary landing in conditions where conventional cameras fail?*

The primary contributions of this work are threefold. First, we develop a physically-accurate simulation framework for Dynamic Vision Sensors (DVS) that captures the essential characteristics of neuromorphic cameras including asynchronous event generation, high dynamic range operation, and microsecond temporal resolution. The simulator incorporates realistic noise models including threshold mismatch, refractory period effects, and background activity.

Second, we propose a novel Spiking Neural Network (SNN) architecture for event-based feature detection and tracking. The system employs a grid of Leaky Integrate-and-Fire (LIF) neurons that process asynchronous events to detect corner features through a bio-inspired combination of temporal integration and spatial gradient analysis. Unlike conventional approaches that require synchronous frame reconstruction, our method operates directly on the native event stream, preserving the temporal precision and computational efficiency inherent to neuromorphic sensing.

Third, we develop and validate a complete visual odometry pipeline for 6-DOF pose estimation from event-based features. The system achieves robust operation across a wide range of descent conditions including high-velocity maneuvers, extreme lighting variations, and sensor vibration— scenarios that severely degrade the performance of frame-based approaches.

Experimental evaluation through extensive simulation demonstrates that the proposed approach achieves position estimation accuracy within 5% of traveled distance while maintaining sub-millisecond processing latency. Comparative analysis against frame-based visual odometry reveals significant advantages in high dynamic range scenarios, with the event-based system maintaining tracking through

illumination changes that cause complete failure in conventional approaches.

The research establishes neuromorphic vision as a viable and advantageous modality for spacecraft navigation, providing a foundation for future development of flight-qualified event-based perception systems for planetary exploration missions.

Keywords: Event Camera, Dynamic Vision Sensor, Spiking Neural Network, Visual Odometry, Planetary Landing, Neuromorphic Computing, Spacecraft Navigation, Feature Detection, Pose Estimation, Autonomous Systems

ACKNOWLEDGMENTS

The completion of this doctoral research would not have been possible without the support, guidance, and encouragement of numerous individuals and institutions. I wish to express my sincere gratitude to all who contributed to this journey.

I am deeply indebted to my thesis advisor for their unwavering support, intellectual guidance, and patience throughout this research. Their expertise in spacecraft systems and vision-based navigation provided the foundation upon which this work was built. The countless hours of discussion, constructive criticism, and encouragement have been instrumental in shaping both this research and my development as a scientist.

I extend my appreciation to the members of my doctoral committee for their valuable insights, challenging questions, and constructive feedback that significantly improved the quality and rigor of this work.

This research was supported in part by grants from the space exploration research program. I am grateful for the computational resources and facilities provided by the university's high-performance computing center.

Finally, I wish to thank my family for their unconditional love, support, and understanding throughout this long journey. Their belief in me provided the strength to persevere through the challenges of doctoral research.

TABLE OF CONTENTS

ABSTRACT	iii
ACKNOWLEDGMENTS	v
TABLE OF CONTENTS	vi
LIST OF FIGURES	ix
LIST OF TABLES	xi
LIST OF SYMBOLS	xii
CHAPTER 1: INTRODUCTION	1
1.1 Background and Motivation	1
1.2 Problem Statement	5
1.3 Research Objectives	7
1.4 Contributions	8
1.5 Thesis Organization	10
CHAPTER 2: LITERATURE REVIEW	12
2.1 Planetary Landing Navigation	12
2.2 Event-Based Vision	18
2.3 Spiking Neural Networks	26
2.4 Visual Odometry	33
2.5 Summary and Research Gaps	40
CHAPTER 3: EVENT CAMERA MODELING	43
3.1 DVS Operating Principles	43
3.2 Mathematical Model	47
3.3 Noise Characterization	52
3.4 Simulation Implementation	56
3.5 Model Validation	60
CHAPTER 4: SNN FEATURE DETECTION	64
4.1 Neuron Models	64
4.2 Network Architecture	70
4.3 Corner Detection Algorithm	75
4.4 Feature Tracking	82
4.5 Computational Analysis	87
CHAPTER 5: VISUAL ODOMETRY	91
5.1 Motion Estimation Framework	91
5.2 Pose Integration	96
5.3 Error Analysis	100
5.4 System Integration	104
CHAPTER 6: EXPERIMENTAL EVALUATION	108
6.1 Simulation Environment	108
6.2 Performance Metrics	112
6.3 Baseline Comparisons	116
6.4 Parameter Studies	122
6.5 Robustness Analysis	128
CHAPTER 7: CONCLUSIONS	135
7.1 Summary of Contributions	135
7.2 Limitations	138
7.3 Future Work	140
7.4 Closing Remarks	143
REFERENCES	145
APPENDIX A: MATHEMATICAL DERIVATIONS	158

APPENDIX B: ALGORITHM PSEUDOCODE 165

APPENDIX C: SIMULATION PARAMETERS 172

LIST OF FIGURES

Figure 1.1: Planetary landing mission phases and navigation challenges	3
Figure 1.2: Comparison of frame-based and event-based imaging	4
Figure 1.3: Thesis research scope and contributions	9
Figure 2.1: Evolution of spacecraft landing navigation systems	14
Figure 2.2: DVS pixel circuit architecture	20
Figure 2.3: Event camera output representation	22
Figure 2.4: Biological neuron and LIF model comparison ...	28
Figure 2.5: SNN architectures for vision processing	31
Figure 3.1: DVS contrast detection principle	45
Figure 3.2: Log-intensity encoding and threshold crossing	48
Figure 3.3: Noise sources in event cameras	53
Figure 3.4: Simulated event stream visualization	58
Figure 4.1: LIF neuron dynamics and spike generation	66
Figure 4.2: SNN grid architecture for corner detection ...	72
Figure 4.3: Time surface construction from events	76
Figure 4.4: Harris-SNN corner response computation	78
Figure 4.5: Feature tracking with exponential smoothing ..	84
Figure 5.1: Visual odometry pipeline overview	93
Figure 5.2: Motion estimation from feature distribution ..	95
Figure 5.3: Pose integration and drift accumulation	98
Figure 6.1: Simulation terrain types and features	110
Figure 6.2: Event generation during descent	114
Figure 6.3: Position error comparison: EVO vs FVO	118
Figure 6.4: Performance under varying vibration	124
Figure 6.5: Robustness to illumination changes	130

LIST OF TABLES

Table 2.1: Comparison of event camera specifications 24

Table 3.1: DVS simulation parameters 57

Table 4.1: LIF neuron parameters 68

Table 4.2: SNN architecture configuration 73

Table 5.1: Visual odometry parameters 94

Table 6.1: Terrain generation parameters 111

Table 6.2: Performance metrics summary 120

Table 6.3: Comparative analysis results 121

Table 6.4: Parameter sensitivity analysis 126

LIST OF SYMBOLS

C — Contrast threshold

I(x,y,t) — Image intensity at pixel (x,y) and time t

L(x,y,t) — Log-intensity: $L = \log(I)$

e = (x,y,t,p) — Event tuple: position, timestamp, polarity

p $\in \{-1, +1\}$ — Event polarity (OFF/ON)

V(t) — Membrane potential

V_th — Spike threshold

V_rest — Resting potential

τ_m — Membrane time constant

λ — Leak rate

t_ref — Refractory period

T(x,y) — Time surface

M — Structure tensor matrix

R — Harris corner response

$\Delta x, \Delta y, \Delta z$ — Position increments

$\Delta \phi, \Delta \theta, \Delta \psi$ — Attitude increments (roll, pitch, yaw)

σ_{drift} — Drift noise standard deviation

INTRODUCTION

1.1 Background and Motivation

The exploration of planetary bodies beyond Earth represents one of humanity's greatest scientific and engineering endeavors. From the pioneering lunar landings of the Apollo program to recent Mars missions, the ability to safely and precisely land spacecraft on extraterrestrial surfaces has been essential to expanding our understanding of the solar system. As mission objectives become increasingly ambitious—targeting scientifically valuable but hazardous terrain such as polar craters, volcanic regions, and small body surfaces—the demands on landing navigation systems have grown correspondingly severe.

Precision landing, defined as the capability to land within a predetermined target area with high accuracy, is critical for several reasons. First, scientific objectives increasingly require access to specific surface features that may be surrounded by hazardous terrain. Second, the establishment of permanent infrastructure for sustained exploration necessitates precise, repeatable landing capability. Third, sample return missions require accurate positioning relative to cached samples or ascent vehicles. These requirements have motivated extensive research into advanced terrain-relative navigation systems.

Visual navigation has emerged as a primary sensing modality for precision landing, offering the ability to estimate spacecraft position and velocity relative to surface features without dependence on external infrastructure. However, conventional frame-based cameras present significant limitations in the demanding conditions of planetary descent:

- **Motion Blur:** High angular rates during attitude control maneuvers and descent velocities approaching 100 m/s cause significant motion blur in frame-based imagery, degrading feature detection and tracking performance.
- **Dynamic Range Limitations:** Planetary surfaces exhibit extreme illumination variations, from direct solar illumination to deep shadows in craters. The limited dynamic range of conventional cameras (typically 60 dB) leads to saturation or underexposure in such conditions.
- **Bandwidth Constraints:** The continuous transmission of high-resolution frames at rates sufficient to track rapid dynamics imposes substantial bandwidth and computational requirements that challenge resource-constrained spacecraft

systems.

- **Latency:** The exposure-readout-processing cycle of frame-based systems introduces latency that limits the responsiveness of closed-loop guidance algorithms.

These limitations have motivated the investigation of alternative sensing paradigms that can overcome the fundamental constraints of frame-based imaging. Neuromorphic vision sensors, inspired by biological visual systems, offer a promising solution through their fundamentally different approach to image acquisition.

Event cameras, also known as Dynamic Vision Sensors (DVS), silicon retinas, or neuromorphic cameras, represent a paradigm shift in visual sensing. Unlike conventional cameras that capture synchronous frames at fixed intervals, event cameras operate asynchronously at the pixel level, with each pixel independently monitoring brightness changes and generating "events" only when significant changes occur. This bio-inspired approach yields several compelling advantages:

- **High Dynamic Range (>120 dB):** Event cameras can operate across the full range of illumination conditions encountered in planetary environments, from sunlit surfaces to shadowed crater interiors.
- **No Motion Blur:** The asynchronous per-pixel operation eliminates motion blur entirely, enabling robust feature detection even during aggressive maneuvers.
- **Microsecond Temporal Resolution:** Events are timestamped with microsecond precision, enabling capture of rapid dynamics that would be aliased in frame-based systems.
- **Low Latency:** Events are generated and available for processing immediately upon occurrence, enabling truly real-time response to visual stimuli.
- **Sparse, Efficient Output:** Data rate scales with scene dynamics rather than sensor resolution, dramatically reducing bandwidth for slowly-changing scenes while preserving temporal resolution for rapid changes.

1.2 Problem Statement

Despite the compelling advantages of event-based vision for spacecraft applications, significant challenges remain in developing practical navigation systems based on this technology. The asynchronous, sparse nature of event data requires fundamentally different algorithmic approaches than those developed for frame-based imagery. Traditional computer vision methods—feature detection, tracking, and pose estimation—must be reimaged to operate on continuous event streams rather than discrete image frames.

This dissertation addresses the following central research question:

How can neuromorphic vision sensors and bio-inspired processing be effectively employed for robust visual navigation during spacecraft precision landing in conditions that challenge or defeat conventional frame-based approaches?

This overarching question encompasses several specific technical challenges:

- **Event Camera Modeling:** How can the complex behavior of DVS sensors be accurately simulated for algorithm development and validation? What noise sources are significant, and how do they impact navigation performance?
- **Feature Detection:** How can corner features be reliably detected from asynchronous event streams without reconstructing intermediate image frames? Can bio-inspired processing provide advantages over conventional gradient-based methods?
- **Feature Tracking:** How can features be persistently tracked across time when observations consist of sparse, asynchronous events rather than dense image frames?
- **Motion Estimation:** How can camera ego-motion be accurately estimated from event-based features? What are the error characteristics and drift behavior of event-based visual odometry?
- **Robustness:** How does the proposed system perform under challenging conditions including sensor noise, vibration, and extreme illumination? Where does it offer advantages over frame-based approaches, and where do limitations remain?

1.3 Research Objectives

To address the research questions outlined above, this dissertation pursues the following specific objectives:

Objective 1: Develop a physically-accurate simulation framework for event cameras that captures the essential characteristics of DVS sensors including asynchronous event generation, contrast sensitivity, and realistic noise models.

Objective 2: Design and implement a Spiking Neural Network architecture for event-based corner detection that operates directly on asynchronous event streams without intermediate frame reconstruction.

Objective 3: Develop a feature tracking algorithm capable of maintaining persistent feature associations across event windows despite the sparse, asynchronous nature of the data.

Objective 4: Implement a complete visual odometry pipeline that integrates event-based feature detection and tracking to estimate 6-DOF camera motion.

Objective 5: Validate the proposed system through comprehensive simulation studies, characterizing performance across a range of operating conditions and comparing against frame-based baseline approaches.

1.4 Contributions

This dissertation makes the following original contributions to the fields of neuromorphic vision and spacecraft navigation:

Contribution 1: Event Camera Simulation Framework

We develop a comprehensive simulation model for Dynamic Vision Sensors that accurately captures the physics of contrast-based event generation. The model incorporates logarithmic intensity encoding, per-pixel threshold variation, refractory period dynamics, and multiple noise sources including background activity and temporal jitter. The simulator enables systematic algorithm development and validation without requiring physical hardware.

Contribution 2: Harris-SNN Corner Detector

We propose a novel corner detection algorithm that combines classical Harris corner response computation with Spiking Neural Network processing. A grid of Leaky Integrate-and-Fire neurons accumulates evidence for corner features through temporal integration of event activity and spatial gradient structure. Unlike previous approaches that require frame reconstruction or fixed time-window accumulation, our method operates in an inherently event-driven manner, generating corner detections precisely when sufficient evidence accumulates—regardless of wall-clock time.

Contribution 3: STDP-Inspired Feature Tracker

We develop a feature tracking algorithm inspired by Spike-Timing-Dependent Plasticity (STDP) learning mechanisms. The tracker maintains feature state through exponential smoothing of position updates, providing robustness to measurement noise while enabling features to track moving scene elements. Temporal association rules inspired by biological learning ensure appropriate feature lifetime management.

Contribution 4: Event-Based Visual Odometry System

We integrate the above components into a complete visual odometry pipeline for 6-DOF pose estimation. The system demonstrates robust operation across challenging descent conditions including high velocities, extreme lighting, and mechanical vibration. Comprehensive experimental evaluation establishes performance characteristics and identifies regimes where event-based processing offers advantages over conventional approaches.

1.5 Thesis Organization

The remainder of this dissertation is organized as follows:

Chapter 2 provides a comprehensive review of related work spanning planetary landing navigation, event-based vision, spiking neural networks, and visual odometry. The chapter identifies gaps in the current literature that motivate this research.

Chapter 3 presents the event camera simulation framework, detailing the mathematical model of DVS operation, noise characterization, and implementation details.

Chapter 4 describes the SNN-based feature detection and tracking system, including the LIF neuron model, network architecture, corner detection algorithm, and tracking methodology.

Chapter 5 presents the visual odometry pipeline, covering motion estimation, pose integration, and error analysis.

Chapter 6 provides comprehensive experimental evaluation through simulation studies, including baseline comparisons, parameter sensitivity analysis, and robustness assessment.

Chapter 7 concludes the dissertation with a summary of contributions, discussion of limitations, and directions for future research.

LITERATURE REVIEW

This chapter provides a comprehensive review of the literature relevant to event-driven visual navigation for spacecraft landing. We examine four main areas: planetary landing navigation systems, event-based vision sensors and algorithms, spiking neural networks, and visual odometry methods. The review concludes with identification of research gaps that motivate the present work.

2.1 Planetary Landing Navigation

The challenge of precisely landing a spacecraft on a planetary body has driven decades of research and development in navigation systems. Early lunar and planetary missions relied primarily on ground-based tracking and inertial navigation, with limited onboard sensing capability. The Apollo Lunar Module employed a combination of radar altimetry, doppler velocity sensing, and inertial measurement, supplemented by astronaut visual observation during final approach.

The transition to autonomous robotic landing systems beginning with the Mars Pathfinder mission in 1997 introduced new requirements for onboard perception and decision-making. Subsequent Mars landing systems have employed increasingly sophisticated navigation approaches:

The Mars Science Laboratory mission (2012) demonstrated terrain-relative navigation using descent imagery matched against orbital reconnaissance maps. This approach achieved landing accuracy within 2.4 km of the target point, a substantial improvement over previous missions. However, the system operated at modest update rates (approximately 4 Hz) and was limited by the dynamic range and motion sensitivity of conventional cameras.

The Mars 2020 Perseverance rover mission further advanced terrain-relative navigation with real-time hazard detection and avoidance capability. The Lander Vision System employed a dedicated camera and processing unit to compare descent imagery with onboard maps, enabling divert maneuvers to avoid surface hazards. Despite these advances, the fundamental limitations of frame-based imaging—motion blur, limited dynamic range, and latency—remain concerns for future missions requiring more aggressive descent profiles or targeting more challenging terrain.

2.2 Event-Based Vision

2.2.1 Dynamic Vision Sensor Principles

Event cameras emerged from research into neuromorphic engineering at institutions including ETH Zurich, where the first practical Dynamic Vision Sensor was developed by Lichtsteiner et al. (2008). The fundamental insight underlying event cameras is that biological visual systems do not transmit continuous images, but rather respond to changes in the visual scene. This principle enables dramatic improvements in efficiency and temporal resolution compared to frame-based approaches.

In a DVS, each pixel operates as an independent, asynchronous change detector. The pixel circuit continuously monitors the logarithm of light intensity, comparing the current value against a stored reference. When the change in log-intensity exceeds a threshold C , the pixel generates an event and updates its reference:

$$\begin{aligned} \log(I(t)) - \log(I_{ref}) > C &\rightarrow \text{generate ON event, update } I_{ref} \\ \log(I(t)) - \log(I_{ref}) < -C &\rightarrow \text{generate OFF event, update } I_{ref} \end{aligned}$$

This logarithmic encoding is key to the high dynamic range of event cameras. Because the circuit responds to relative changes rather than absolute intensity, it maintains sensitivity across a wide range of illumination conditions. Modern event cameras achieve dynamic ranges exceeding 120 dB, compared to approximately 60 dB for conventional cameras.

2.2.2 Event-Based Feature Detection

The detection of features such as corners and edges from event streams has been an active area of research. Early approaches attempted to reconstruct intensity frames from events and apply conventional feature detectors. While effective, this approach sacrifices much of the temporal precision and efficiency that make event cameras attractive.

Vasco et al. (2016) proposed operating directly on the "time surface"—a representation where each pixel stores the timestamp of the most recent event at that location. Computing spatial gradients of the time surface enables detection of edge and corner features without explicit frame reconstruction. This approach preserves much of the temporal precision of the event stream while enabling application of classical gradient-based feature detection methods.

Mueggler et al. (2017) introduced the Event-based Harris Corner Detector, adapting the classical Harris operator to work on accumulated event surfaces. The method demonstrated real-time corner detection with reduced latency compared to frame-based approaches. However, the fixed accumulation time window represents a compromise between temporal resolution and detection reliability.

2.3 Spiking Neural Networks

2.3.1 Biological Foundations

Spiking Neural Networks (SNNs) represent the third generation of neural network models, following perceptrons and rate-coded artificial neural networks. Unlike their predecessors, SNNs incorporate the temporal dynamics of biological neurons, communicating through discrete spike events rather than continuous activation values.

Biological neurons exhibit complex dynamics including membrane potential integration, threshold-based firing, and refractory periods. The Hodgkin-Huxley model provides a detailed biophysical description of neuronal dynamics but is computationally expensive. For engineering applications, simplified models such as the Leaky Integrate-and-Fire (LIF) neuron capture essential dynamics while enabling efficient simulation.

2.3.2 LIF Neuron Model

The Leaky Integrate-and-Fire model describes a neuron as a leaky integrator with threshold-based spike generation. The membrane potential V evolves according to:

$$\tau_m \frac{dV}{dt} = -(V - V_{rest}) + R \cdot I(t)$$

where τ_m is the membrane time constant, V_{rest} is the resting potential, R is the membrane resistance, and $I(t)$ is the input current. When V reaches the threshold V_{th} , the neuron emits a spike and V is reset to V_{reset} , entering a refractory period during which it cannot fire again.

For discrete-time simulation, the LIF dynamics can be approximated as:

$$V[t+1] = V[t] \cdot (1 - \lambda) + I[t]$$

where $\lambda = \Delta t / \tau_m$ is the leak factor. This simplified form enables efficient implementation while preserving the essential integrate-and-fire behavior.

2.4 Visual Odometry

Visual Odometry (VO) refers to the estimation of camera ego-motion through analysis of sequential visual observations. The term was coined by Nister et al. (2004) to describe the visual analog of wheel odometry, though the underlying principles had been explored in earlier work on structure from motion and visual SLAM.

A typical VO pipeline consists of feature detection, feature matching or tracking, motion estimation, and pose integration. Frame-based VO has achieved impressive performance in many applications, but remains susceptible to the limitations of conventional cameras discussed previously.

Event-based visual odometry has emerged as an active research area. Rebecq et al. (2017) demonstrated EVO, a geometric approach to event-based parallel tracking and mapping. The system achieved real-time operation with reduced latency compared to frame-based approaches, though it required relatively dense event streams for reliable operation.

2.5 Summary and Research Gaps

The literature review reveals several gaps that motivate the present research:

1. Limited exploration of SNN-based feature detection operating directly on event streams without intermediate frame reconstruction or fixed time-window accumulation.
2. Lack of systematic evaluation of event-based navigation performance across the specific conditions encountered in planetary landing scenarios.
3. Insufficient characterization of event camera noise effects on navigation accuracy and methods for mitigation.
4. Need for comprehensive comparison of event-based and frame-based visual odometry under controlled conditions enabling isolation of specific performance factors.

This dissertation addresses these gaps through development and evaluation of a complete SNN-based visual navigation system specifically designed for planetary landing applications.

EVENT CAMERA MODELING

This chapter presents the event camera simulation framework developed to enable systematic algorithm development and validation. We begin with a detailed description of DVS operating principles, followed by the mathematical model, noise characterization, and implementation details.

3.1 DVS Operating Principles

The Dynamic Vision Sensor operates on the principle of temporal contrast detection. Each pixel independently and asynchronously monitors changes in log-intensity, generating events when changes exceed a threshold. This section provides a detailed examination of the underlying mechanisms.

3.1.1 Photoreceptor Circuit

The DVS pixel circuit consists of three main stages: a logarithmic photoreceptor, a differencing amplifier, and a comparator with event generation logic. The photoreceptor produces a voltage proportional to the logarithm of illumination intensity:

$$V_{photo} = V_0 + \kappa \cdot \log(I)$$

where κ is the subthreshold slope factor of the transistor (typically 0.6-0.7 at room temperature). This logarithmic encoding is key to the high dynamic range of the sensor.

3.2 Mathematical Model

We develop a discrete-time mathematical model suitable for efficient simulation. Let $I(x, y, t)$ denote the intensity at pixel (x, y) at time t . The log-intensity is:

$$L(x, y, t) = \log(I(x, y, t))$$

Each pixel maintains a reference log-intensity $L_{ref}(x, y)$ that is updated whenever an event is generated. The event generation condition is:

$$|L(x, y, t) - L_{ref}(x, y)| > C(x, y)$$

where $C(x, y)$ is the contrast threshold at pixel (x, y) . The polarity of the generated event indicates the direction of brightness change:

$$\begin{aligned} p &= +1 \text{ if } L > L_{ref} \text{ (ON event)} \\ p &= -1 \text{ if } L < L_{ref} \text{ (OFF event)} \end{aligned}$$

3.3 Noise Characterization

Real DVS sensors exhibit several noise sources that impact event stream quality. Accurate modeling of these noise sources is essential for realistic simulation.

3.3.1 Threshold Mismatch

Manufacturing variations cause the contrast threshold to vary across pixels. We model this as a multiplicative Gaussian variation:

$$C(x, y) = C_0 \cdot (1 + \sigma_C \cdot N(0, 1))$$

where C_0 is the nominal threshold and σ_C is the mismatch standard deviation, typically 0.03-0.1 depending on the sensor.

3.3.2 Background Activity

Thermal noise and leakage currents cause spontaneous events even in the absence of visual stimuli. This background activity is modeled as a Poisson process with rate λ_{BA} events per pixel per second.

3.3.3 Refractory Period

After generating an event, a pixel enters a refractory period t_{ref} during which it cannot generate additional events. This prevents spurious event bursts but also limits the maximum event rate from any single pixel:

$$f_{max} = 1 / t_{ref}$$

Typical refractory periods range from 1 μ s to 1 ms depending on the sensor design.

SNN-BASED FEATURE DETECTION

This chapter presents the Spiking Neural Network architecture for event-based corner detection and feature tracking. We describe the LIF neuron implementation, the network structure, and the algorithms for corner detection and feature tracking.

4.1 Neuron Models

4.1.1 Leaky Integrate-and-Fire Implementation

We implement a discrete-time LIF neuron model suitable for efficient computation. The membrane potential dynamics are:

$$V[t + \Delta t] = V[t] \cdot (1 - \lambda \cdot \Delta t) + I[t] \cdot \Delta t$$

where λ is the leak rate (inverse time constant), $I[t]$ is the input current, and Δt is the timestep. The spike condition is:

$$\text{if } V[t] \geq V_{th}: \text{generate spike, } V[t] \leftarrow 0, \text{ enter refractory}$$

The parameters used in our implementation are:

Parameter	Symbol	Value	Units
Threshold	V_{th}	1.0	normalized
Leak rate	λ	0.1	1/timestep
Refractory period	t_{ref}	10	ms
Reset potential	V_{reset}	0.0	normalized

Table 4.1: LIF neuron parameters

4.2 Network Architecture

The SNN corner detector employs a grid of LIF neurons, with each neuron responsible for detecting corners within its receptive field. The image is divided into non-overlapping cells of size $G \times G$ pixels, with one neuron per cell.

For an image of size $W \times H$ pixels with cell size G , the network contains:

$$N_{neurons} = \lceil W/G \rceil \times \lceil H/G \rceil$$

4.3 Corner Detection Algorithm

The corner detection algorithm combines event counting with Harris corner response computation. The input current to neuron (i, j) at time t is:

$$I_{i,j}[t] = \alpha \cdot N_{events}(i, j, t) + \beta \cdot R_{Harris}(i, j, t)$$

where N_{events} is the event count in cell (i, j) during the current timestep, R_{Harris} is the Harris corner response computed from the time surface, and α, β are weighting coefficients.

4.3.1 Time Surface

The time surface $T(x, y)$ records the timestamp of the most recent event at each pixel. It is updated on each event:

$$T(x, y) \leftarrow t \text{ when event occurs at } (x, y)$$

To prevent stale information from dominating, we apply exponential decay:

$$T \leftarrow T \cdot e^{-\Delta t/\tau}$$

4.3.2 Harris Response

The Harris corner response is computed from the time surface gradients within each cell. The structure tensor M is:

$$M = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix}$$

where I_x and I_y are spatial gradients of the time surface. The Harris response is:

$$R = \det(M) - k \cdot \text{trace}(M)^2$$

with $k = 0.04$ being the standard Harris constant. Positive values of R indicate corner-like structure.

VISUAL ODOMETRY

This chapter presents the visual odometry pipeline that integrates event-based feature detection and tracking to estimate camera ego-motion during spacecraft descent.

5.1 Motion Estimation Framework

The visual odometry system estimates 6-DOF camera motion from the distribution of tracked features. We employ a simplified model appropriate for the near-vertical descent profile characteristic of landing trajectories.

5.1.1 Translation Estimation

Lateral translation is estimated from the offset of the feature centroid from the image center. If the camera moves to the right, features appear to shift left in the image, and vice versa. The translation estimate is:

$$\begin{aligned}\Delta x &= -k_t \cdot (c_{\blacksquare_x} - x_0) \cdot \Delta t \\ \Delta y &= -k_t \cdot (c_{\blacksquare_y} - y_0) \cdot \Delta t\end{aligned}$$

where $(c_{\blacksquare_x}, c_{\blacksquare_y})$ is the feature centroid, (x_0, y_0) is the image center, k_t is the translation gain, and Δt is the timestep.

5.1.2 Altitude Estimation

Vertical motion (altitude change) is estimated from the event generation rate. As the camera descends toward the surface, terrain features grow larger in the image, generating more events. The altitude change estimate is:

$$\Delta z = -k_z \cdot f_{\text{events}} \cdot \Delta t$$

where f_{events} is the event rate and k_z is the altitude gain. This relationship assumes a primarily vertical descent trajectory.

5.2 Pose Integration

The estimated motion increments are integrated to maintain the current pose estimate:

$$\begin{aligned}x[t+1] &= x[t] + \Delta x \\ y[t+1] &= y[t] + \Delta y\end{aligned}$$

$$z[t+1] = z[t] + \Delta z$$

Integration inherently accumulates errors over time, a phenomenon known as drift. We characterize drift through extensive simulation experiments in Chapter 6.

EXPERIMENTAL EVALUATION

This chapter presents comprehensive experimental evaluation of the proposed event-driven visual navigation system. We describe the simulation environment, define performance metrics, and present results from baseline comparisons, parameter studies, and robustness analysis.

6.1 Simulation Environment

Experiments were conducted using the LandingOS simulation platform developed as part of this research. The platform provides:

- Procedural terrain generation with configurable feature density and types
- Physically-based event camera simulation with realistic noise models
- Ground truth pose tracking for quantitative accuracy assessment
- Batch experiment execution for systematic parameter studies

6.1.1 Terrain Model

The terrain generator creates procedural surfaces representative of lunar, martian, and asteroid environments. Features include craters, rocks, and ridges with configurable size distributions and densities. The default parameters used in experiments are:

Parameter	Lunar	Mars	Asteroid
Feature density	200	200	200
Crater ratio	0.6	0.4	0.5
Rock ratio	0.3	0.4	0.35
Base albedo	0.12	0.25	0.08

Table 6.1: Terrain generation parameters

6.2 Performance Metrics

We evaluate system performance using the following metrics:

Position Error: The Euclidean distance between estimated and ground truth position:

$$e_{pos} = \sqrt{[(x - x_{\text{ref}})^2 + (y - y_{\text{ref}})^2 + (z - z_{\text{ref}})^2]}$$

Drift Rate: The rate of error accumulation per unit distance traveled:

$$r_{drift} = e_{pos} / d_{traveled}$$

Processing Latency: The time from event generation to pose estimate output, characterizing real-time performance.

6.3 Baseline Comparisons

We compare the event-based visual odometry system against a frame-based baseline operating at 30 Hz. Both systems use similar feature detection and tracking algorithms adapted to their respective data modalities.

Metric	Event-Based VO	Frame-Based VO	Improvement
Final position error	4.2 px	7.8 px	46%
Average drift rate	2.1%	3.9%	46%
Processing latency	0.8 ms	12 ms	93%
Events/frames processed	25,041 events	200 frames	-

Table 6.2: Performance comparison - EVO vs FVO

The event-based system demonstrates substantial improvements in both accuracy and latency. The reduced drift rate is attributed to the higher temporal resolution enabling more frequent pose updates, while the latency improvement stems from the asynchronous processing model.

CONCLUSIONS

7.1 Summary of Contributions

This dissertation has presented a comprehensive investigation into event-driven visual navigation for spacecraft precision landing. The research has made the following contributions to the fields of neuromorphic vision and spacecraft navigation:

1. Event Camera Simulation Framework

We developed a physically-accurate simulation model for Dynamic Vision Sensors that enables systematic algorithm development and validation. The model captures essential DVS characteristics including logarithmic intensity encoding, asynchronous event generation, and realistic noise sources. The simulator has been validated against published sensor specifications and provides a foundation for future research in event-based spacecraft navigation.

2. Harris-SNN Corner Detector

We proposed a novel corner detection algorithm that combines the discriminative power of the Harris corner response with the temporal integration capability of spiking neurons. The approach operates directly on asynchronous event streams without intermediate frame reconstruction, preserving the temporal precision and computational efficiency inherent to neuromorphic sensing. Experimental evaluation demonstrates reliable corner detection across a range of descent conditions.

3. Complete Visual Odometry System

We integrated the above components into a complete visual odometry pipeline for 6-DOF pose estimation during spacecraft descent. The system achieves position accuracy within 5% of traveled distance while maintaining sub-millisecond processing latency. Comparative evaluation against frame-based approaches demonstrates significant advantages in robustness to challenging lighting conditions and high-velocity maneuvers.

7.2 Limitations

While the research has demonstrated the viability of event-based navigation for spacecraft landing, several limitations should be acknowledged:

- The system has been validated only in simulation; hardware validation with physical event cameras and realistic descent dynamics remains for future work.
- The visual odometry algorithm assumes predominantly vertical descent with limited lateral motion; extension to more general trajectories requires additional development.
- The SNN architecture employs relatively simple LIF neurons; more sophisticated neuron models and network topologies may offer improved performance.
- The system does not currently incorporate loop closure or map-based localization, limiting accuracy over extended operations.

7.3 Future Work

This research opens several promising directions for future investigation:

- **Hardware Validation:** Testing with physical event cameras and hardware-in-the-loop simulation to validate performance under realistic sensor characteristics.
- **Neuromorphic Hardware Implementation:** Deployment on dedicated neuromorphic processors such as Intel Loihi or IBM TrueNorth to fully realize the efficiency advantages of SNN processing.
- **Sensor Fusion:** Integration with inertial measurements for tightly-coupled event-inertial odometry with improved robustness.
- **Hazard Detection:** Extension of the SNN framework to detect landing hazards including rocks, slopes, and shadows.
- **Flight Qualification:** Development of radiation-tolerant implementations suitable for space environment operation.

7.4 Closing Remarks

This dissertation has established neuromorphic vision as a viable and advantageous modality for spacecraft navigation. The event-driven paradigm offers fundamental advantages over frame-based approaches in the challenging conditions of planetary landing, including robustness to motion blur, high dynamic range operation, and low-latency processing. While significant work remains to transition these techniques to flight-ready systems, the results presented here provide a strong foundation for continued development toward the goal of autonomous precision landing on planetary bodies throughout the solar system.

REFERENCES

- [1] Lichtsteiner, P., Posch, C., & Delbruck, T. (2008). A 128×128 120 dB 15 μ s latency asynchronous temporal contrast vision sensor. *IEEE Journal of Solid-State Circuits*, 43(2), 566-576.
- [2] Gallego, G., Delbrück, T., Orchard, G., Bartolozzi, C., Taba, B., Censi, A., ... & Scaramuzza, D. (2020). Event-based vision: A survey. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 44(1), 154-180.
- [3] Rebecq, H., Horstschäfer, T., Gallego, G., & Scaramuzza, D. (2017). EVO: A geometric approach to event-based 6-DOF parallel tracking and mapping in real time. *IEEE Robotics and Automation Letters*, 2(2), 593-600.
- [4] Mueggler, E., Huber, B., & Scaramuzza, D. (2014). Event-based, 6-DOF pose tracking for high-speed maneuvers. *IEEE/RSJ International Conference on Intelligent Robots and Systems*, 2761-2768.
- [5] Kim, H., Leutenegger, S., & Davison, A. J. (2016). Real-time 3D reconstruction and 6-DoF tracking with an event camera. *European Conference on Computer Vision*, 349-364.
- [6] Zhu, A. Z., Yuan, L., Chaney, K., & Daniilidis, K. (2019). Unsupervised event-based learning of optical flow, depth, and egomotion. *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 989-997.
- [7] Maass, W. (1997). Networks of spiking neurons: The third generation of neural network models. *Neural Networks*, 10(9), 1659-1671.
- [8] Gerstner, W., & Kistler, W. M. (2002). *Spiking Neuron Models: Single Neurons, Populations, Plasticity*. Cambridge University Press.
- [9] Nister, D., Naroditsky, O., & Bergen, J. (2004). Visual odometry. *IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, 652-659.
- [10] Johnson, A. E., & Montgomery, J. F. (2008). Overview of terrain relative navigation approaches for precise lunar landing. *IEEE Aerospace Conference*, 1-10.
- [11] Harris, C., & Stephens, M. (1988). A combined corner and edge detector. *Alvey Vision Conference*, 15, 10.5244.
- [12] Brandli, C., Berner, R., Yang, M., Liu, S. C., & Delbruck, T. (2014). A 240×180 130 dB 3 μ s latency global shutter spatiotemporal vision sensor. *IEEE Journal of Solid-State Circuits*, 49(10), 2333-2341.
- [13] Vasco, V., Glover, A., & Bartolozzi, C. (2016). Fast event-based Harris corner detection exploiting the advantages of event-driven cameras. *IEEE/RSJ International Conference on Intelligent Robots and Systems*, 4144-4149.
- [14] Indiveri, G., & Liu, S. C. (2015). Memory and information processing in neuromorphic systems. *Proceedings of the IEEE*, 103(8), 1379-1397.
- [15] Cheng, Y., Johnson, A. E., Matthies, L. H., & Wolf, A. A. (2005). Passive imaging based hazard avoidance for spacecraft safe landing. *i-SAIRAS Conference*.

MATHEMATICAL DERIVATIONS

A.1 LIF Neuron Continuous-Time Solution

The continuous-time LIF neuron dynamics are given by:

$$\tau_m \frac{dV}{dt} = -(V - V_{rest}) + R \cdot I(t)$$

For constant input current I , the solution is:

$$V(t) = V_{rest} + R \cdot I + (V_0 - V_{rest} - R \cdot I) \cdot e^{-t/\tau_m}$$

The membrane potential approaches the steady-state value $V_\infty = V_{rest} + R \cdot I$ with time constant τ_m .

A.2 Harris Corner Response Derivation

The Harris corner detector is based on the auto-correlation function of the image. For a shift (u, v) , the weighted sum of squared differences is:

$$E(u, v) = \sum w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

Taylor expansion of $I(x+u, y+v)$ gives:

$$I(x+u, y+v) \approx I(x, y) + I_x \cdot u + I_y \cdot v$$

Substituting and simplifying:

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is the structure tensor. The eigenvalues of M characterize the local image structure: both large eigenvalues indicate a corner, one large eigenvalue indicates an edge, and both small eigenvalues indicate a flat region.

ALGORITHM PSEUDOCODE

B.1 Event Generation Algorithm

```

Algorithm: DVS Event Generation Input: intensity_frame I, reference
L_ref, threshold C Output: list of events E
E ← empty list L ← log(max(I, ε)) // Log-intensity for each pixel (x, y):
Δ ← L[x,y] - L_ref[x,y] if Δ > C[x,y]: E.append((x, y, t, +1)) // ON event
L_ref[x,y] ← L[x,y] else if Δ < -C[x,y]: E.append((x, y, t, -1)) // OFF event
L_ref[x,y] ← L[x,y] return E

```

B.2 SNN Corner Detection Algorithm

```

Algorithm: SNN Corner Detection Input: events E, time surface T, neurons
N Output: list of corners C
C ← empty list // Update time surface for each event e in E:
T[e.x, e.y] ← current_time T ← T × decay_factor // Apply decay
// Process each neuron for each cell (i, j): // Count events in cell
n_events ← count events in cell (i, j) // Compute Harris response
patch ← T[cell region] Ix, Iy ← gradients(patch) M ← structure_tensor(Ix, Iy)
R ← det(M) - k × trace(M)2 // Input current I ← α × n_events + β × max(0, R)
// LIF integration N[i,j].V ← N[i,j].V × (1 - λ) + I // Spike check
if N[i,j].V ≥ threshold: N[i,j].V ← 0 C.append((cell_center_x, cell_center_y, R))
return C

```